IEEE Access
Multidisciplinary ⋮ Rapid Review ⋮ Open Access Journal

Received 13 December 2024, accepted 2 January 2025, date of publication 6 January 2025, date of current version 15 January 2025.

Digital Object Identifier 10.1109/ACCESS.2025.3526239

RESEARCH ARTICLE

GNN-EADD: Graph Neural Network-Based E-Commerce Anomaly Detection via Dual-Stage Learning

ZHOUHANG SHAO[1], XURAN WANG[2], ENKAI JI[3], SHIYANG CHEN[4], (Member, IEEE), AND JIN WANG[5]

[1]Department of Computer Science and Engineering, University of California San Diego, La Jolla, CA 92093, USA
[2]Department of Computer and Information Science, University of Pennsylvania, Philadelphia, PA 19104, USA
[3]Department of Computer Science, Rutgers University, New Brunswick, NJ 08901, USA
[4] College of Engineering, Texas A&M University, College Station, TX 77843, USA
[5]Viterbi School of Engineering, University of Southern California, Los Angeles, CA 90089, USA

Corresponding author: Zhouhang Shao (z2shao@ucsd.edu)

ABSTRACT E-commerce platforms face significant challenges in detecting anomalous products, including counterfeit goods and fraudulent listings, which can undermine user trust and platform integrity. This paper presents Graph Neural Network-based E-commerce Anomaly Detection via Dual-stage Learning (GNN-EADD), a novel approach leveraging graph neural networks for anomaly detection in large-scale e-commerce ecosystems. Our key contributions include: 1) A heterogeneous graph representation incorporating products, sellers, and buyers as nodes with their relationships as edges; 2) A novel dual-stage learning framework combining unsupervised graph embedding with semi-supervised fine-tuning; and 3) An attention mechanism that effectively captures complex patterns within network structures. Extensive experiments on a large-scale Amazon dataset demonstrate that GNN-EADD significantly outperforms state-of-the-art baselines in terms of anomaly detection accuracy, precision, and recall, while showing robustness to various types of anomalies and scalability to large networks.

INDEX TERMS Graph neural networks, e-commerce, anomaly detection, heterogeneous graphs, graph attention networks.

I. INTRODUCTION

The rapid growth of e-commerce platforms has revolutionized the way people shop, offering unprecedented convenience and access to a vast array of products. However, this expansion has also brought significant challenges [1], [2], [3], particularly in maintaining the integrity and trustworthiness of online marketplaces. One of the most pressing issues facing e-commerce platforms is the detection of anomalous products, including counterfeit goods, fraudulent listings, and other forms of malicious activities [4], [5], [6].

The prevalence of such anomalies can have severe consequences for both consumers and legitimate businesses. Counterfeit products not only lead to financial losses for consumers and brand owners but can also pose serious health and safety risks [7]. Fraudulent listings, such as fake reviews or artificially inflated ratings, manipulate consumer trust and distort market dynamics [8]. Consequently, developing robust and efficient methods for detecting these anomalies has become a critical priority for e-commerce platforms to ensure user trust, maintain platform integrity, and foster a healthy online ecosystem [9].

Traditional approaches to anomaly detection in e-commerce have primarily relied on rule-based systems or simple machine learning models that analyze individual product features or user behaviors in isolation [10]. However, these methods often fall short in capturing the complex, interconnected nature of e-commerce ecosystems. They struggle to identify sophisticated fraud schemes that exploit relationships between multiple entities, such as coordinated

The associate editor coordinating the review of this manuscript and approving it for publication was Chao Tong.

© 2025 The Authors. This work is licensed under a Creative Commons Attribution 4.0 License.
For more information, see https://creativecommons.org/licenses/by/4.0/

networks of fake accounts or subtle patterns of review manipulation [11].

Recent advancements in graph-based methods, particularly Graph Neural Networks (GNNs), have shown promising results in various domains by effectively modeling and learning from complex relational data [12]. In the context of e-commerce, the inherent graph structure of marketplaces—where products, sellers, and buyers form a natural heterogeneous network—makes GNNs an appealing approach for anomaly detection [13].

Despite the potential of GNNs, several challenges remain in applying them effectively to e-commerce anomaly detection: 1) Heterogeneity: E-commerce graphs are inherently heterogeneous, containing different types of nodes (products, sellers, buyers) and edges (purchases, reviews, co-purchasing). Existing GNN models often struggle to capture this rich heterogeneity effectively [14]. 2) Scalability: E-commerce platforms often involve millions of products and users, necessitating methods that can scale to massive graphs [15]. 3) Limited Labels: Obtaining ground truth labels for anomalous products is expensive and time-consuming, making fully supervised approaches impractical [16]. 4) Dynamic Nature: E-commerce ecosystems are highly dynamic, with new products, users, and interactions constantly emerging, requiring models that can adapt to evolving patterns [17].

To address these challenges, we propose Graph Neural Network-based E-commerce Anomaly Detection via Dual-stage Learning (GNN-EADD), a novel approach that leverages the power of graph neural networks to identify anomalous products in large-scale e-commerce ecosystems. Our method constructs a comprehensive heterogeneous graph representation that captures the complex relationships between products, sellers, and buyers, enabling the detection of sophisticated fraud patterns that might be missed when considering entities in isolation. To handle the limited availability of labeled anomalies while leveraging abundant unlabeled data, we design a dual-stage learning framework. In the first stage, a graph autoencoder learns initial node embeddings in an unsupervised manner to capture the inherent structure of the e-commerce graph. The second stage employs a type-specific Graph ATtention network (GAT) to fine-tune these embeddings using limited labeled anomalies, providing interpretable insights into the detection process. Our approach also incorporates scalable techniques to efficiently handle large-scale e-commerce graphs and adapt to evolving patterns.

The key contributions of this work are:

1) A novel heterogeneous graph representation that effectively captures multiple types of interactions in e-commerce platforms, enabling comprehensive anomaly detection.
2) A dual-stage learning framework combining unsupervised graph embedding with semi-supervised fine-tuning, effectively addressing the challenge of limited labeled anomalies.
3) A type-specific attention mechanism that provides interpretable insights by highlighting suspicious patterns and relationships in the detection process.
4) Extensive experimental validation on a large-scale Amazon dataset, demonstrating significant improvements over state-of-the-art baselines (8.9% in accuracy, 12.5% in precision, and 10.2% in recall) and robust performance across various types of anomalies.

The rest of this paper is organized as follows: Section II reviews related work in graph-based anomaly detection and its applications in e-commerce. Section III provides preliminaries on graph neural networks and anomaly detection. Section IV details our proposed GNN-EADD method, including the heterogeneous graph construction, dual-stage learning framework, and model architecture. Section V presents our experimental setup, datasets, and evaluation metrics. Section VI discusses the results and provides in-depth analysis of our method's performance. Finally, Section 7 concludes the paper and outlines directions for future work.

II. RELATED WORKS

In this section, we review the relevant literature on anomaly detection in e-commerce, graph-based approaches for fraud detection, and the application of graph neural networks to these problems.

A. ANOMALY DETECTION IN E-COMMERCE

Anomaly detection in e-commerce has been a long-standing challenge, with early approaches relying on rule-based systems and traditional machine learning methods. Rayana and Akoglu [10] proposed a method for detecting opinion spam by combining review text, metadata, and relational data. Their work highlighted the importance of considering multiple data sources for effective fraud detection. As e-commerce platforms grew in complexity, researchers began to explore more sophisticated techniques. Li et al. [7] introduced a semi-supervised approach for product classification, which, while not directly addressing anomaly detection, demonstrated the potential of leveraging limited labeled data in e-commerce tasks. Kumar et al. [8] developed REV2, a fraudulent user prediction system for rating platforms, which uses iterative algorithms to compute fairness values of users and reliability of ratings. More recently, Dou et al. [9] addressed the challenge of camouflaged fraudsters in e-commerce, proposing a graph neural network-based approach that enhances fraud detectors against sophisticated fraudulent behaviors. Their work underscores the growing importance of graph-based methods in tackling evolving fraud patterns.

B. GRAPH-BASED ANOMALY DETECTION

The use of graph-based methods for anomaly detection has gained significant traction in recent years. Pourhabibi et al. [4] and Ma et al. [5] provided

comprehensive surveys of graph-based anomaly detection techniques, highlighting their effectiveness in capturing complex relational patterns. In the context of e-commerce, graph-based approaches have shown particular promise. Liu et al. [11] addressed the inconsistency problem in applying graph neural networks to fraud detection, proposing a novel model that better handles the dynamic nature of e-commerce graphs. Their work demonstrated the potential of graph-based methods to capture the evolving relationships between users, products, and transactions in online marketplaces.

C. GRAPH NEURAL NETWORKS FOR ANOMALY DETECTION

The advent of Graph Neural Networks has opened new avenues for anomaly detection in complex, interconnected systems like e-commerce platforms. Wu et al. [12] provided a comprehensive survey of GNNs, outlining their potential applications across various domains, including anomaly detection. Ying et al. [13] demonstrated the scalability of GNNs in web-scale recommender systems, a closely related problem to anomaly detection in e-commerce. Their work showed that GNNs could effectively capture the complex interactions between users and items in large-scale e-commerce platforms. Wang et al. [14] introduced the Heterogeneous Graph Attention Network (HAN), which addresses the challenge of learning on graphs with different types of nodes and edges. This approach is particularly relevant to e-commerce anomaly detection, where products, users, and transactions form a naturally heterogeneous graph.

D. SEMI-SUPERVISED LEARNING IN GRAPH-BASED ANOMALY DETECTION

Given the scarcity of labeled anomalies in real-world e-commerce data, semi-supervised learning approaches have gained prominence. Zhao et al. [16] provided a comprehensive review of graph-based semi-supervised learning methods, highlighting their potential in scenarios with limited labeled data. In the context of e-commerce, semi-supervised methods have shown promise in leveraging both labeled and unlabeled data for improved anomaly detection. Zeng et al. [15] introduced GraphSAINT, an inductive learning method that enables efficient training of GNNs on large-scale graphs, making it particularly suitable for e-commerce applications.

E. DYNAMIC GRAPH LEARNING FOR E-COMMERCE

The dynamic nature of e-commerce platforms, where new products, users, and interactions constantly emerge, poses unique challenges for anomaly detection systems. Pareja et al. [17] proposed EvolveGCN, a GNN model designed for dynamic graphs. While not specifically tailored for e-commerce, their approach offers insights into handling the temporal aspects of e-commerce graphs.

F. RESEARCH GAP AND OUR CONTRIBUTIONS

While the aforementioned works have made significant strides in graph-based anomaly detection for e-commerce, several challenges remain. First, most existing approaches struggle to effectively model the heterogeneous nature of e-commerce graphs, where different types of nodes and edges coexist. Second, the scalability of these methods to large-scale e-commerce platforms with millions of products and users remains a concern. Third, the limited availability of labeled anomalies in real-world scenarios is not adequately addressed by many current approaches.

Our proposed GNN-EADD method addresses these gaps by introducing a dual-stage learning framework that combines unsupervised graph embedding with semi-supervised fine-tuning. This approach allows us to leverage the strengths of both unsupervised learning (for capturing the overall structure of the e-commerce graph) and semi-supervised learning (for fine-tuning with limited labeled anomalies). Furthermore, our method is designed to handle heterogeneous graphs and scale to large e-commerce networks, making it particularly suitable for real-world applications.

III. PRELIMINARIES

In this section, we introduce the fundamental concepts and notations used throughout this paper, including graph theory basics, graph neural networks, and anomaly detection in graphs.

A. GRAPH BASICS

A graph G is defined as a tuple $G = (V, E)$, where $V = \{v_1, v_2, \ldots, v_n\}$ is the set of nodes (or vertices) and $E \subseteq V \times V$ is the set of edges. In the context of e-commerce, nodes can represent entities such as products, users, or sellers, while edges represent relationships or interactions between these entities.

For a heterogeneous graph, which is common in e-commerce scenarios, we extend this definition to include different types of nodes and edges. We define a heterogeneous graph as $G = (V, E, \phi, \psi)$, where:

- V is the set of nodes
- E is the set of edges
- $\phi : V \to A$ is a function that maps each node to its type
- $\psi : E \to R$ is a function that maps each edge to its type
- A is the set of node types (e.g., product, user, seller)
- R is the set of edge types (e.g., purchase, review, sell)

The adjacency matrix $A \in \mathbb{R}^{n \times n}$ of a graph is defined as:

$$A_{ij} = \begin{cases} 1, & \text{if } (v_i, v_j) \in E \\ 0, & \text{otherwise} \end{cases} \tag{1}$$

B. GRAPH NEURAL NETWORKS

GNNs are a class of deep learning models designed to operate on graph-structured data. The key idea behind GNNs is to learn node representations by aggregating information from neighboring nodes.

A general formulation of the node update process in a GNN layer can be expressed as:

$$h_v^{(l+1)} = \sigma\left(W^{(l)} \cdot \text{AGGREGATE}\left(\{h_u^{(l)} : u \in \mathcal{N}(v)\}\right) + b^{(l)}\right) \tag{2}$$

where:

- $h_v^{(l)}$ is the feature vector of node v at layer l
- $\mathcal{N}(v)$ is the set of neighbors of node v
- $W^{(l)}$ and $b^{(l)}$ are learnable parameters
- σ is a non-linear activation function
- AGGREGATE is a function that combines the features of neighboring nodes

C. GRAPH ATTENTION NETWORKS

Graph Attention Networks [18] introduce attention mechanisms to GNNs, allowing the model to assign different importance to different neighbors. The attention coefficient α_{ij} between nodes i and j is computed as:

$$\alpha_{ij} = \frac{\exp(LeakyReLU(a^T[Wh_i \| Wh_j]))}{\sum_{k \in \mathcal{N}(i)} \exp(LeakyReLU(a^T[Wh_i \| Wh_k]))} \tag{3}$$

where a is a learnable attention vector, W is a weight matrix, and $\|$ denotes concatenation.

D. GRAPH AUTOENCODERS

Graph Autoencoders (GAEs) [19] are unsupervised learning models that aim to learn low-dimensional representations of graph nodes. A GAE consists of an encoder that maps nodes to a latent space and a decoder that reconstructs the graph structure from these latent representations.

The encoder of a GAE can be any GNN model. The decoder typically reconstructs the adjacency matrix $\hat{A}$ as:

$$\hat{A} = \sigma(ZZ^T) \tag{4}$$

where Z is the matrix of latent node representations and σ is typically the sigmoid function.

E. ANOMALY DETECTION IN GRAPHS

In the context of graphs, anomalies (or outliers) are nodes, edges, or subgraphs that significantly deviate from the expected pattern of connectivity or attributes in the graph [20]. In e-commerce, these anomalies could represent fraudulent users, fake products, or suspicious transactions.

The anomaly score of a node v can be generally defined as a function f that measures how much v deviates from the normal pattern:

$$\text{anomaly_score}(v) = f(v, G) \tag{5}$$

where G is the graph and f can be based on various criteria such as structural properties, attribute values, or learned representations.

In our work, we focus on semi-supervised anomaly detection, where we have a small set of labeled anomalies to guide the learning process. Let $Y = \{y_1, y_2, \ldots, y_n\}$ be the set of node labels, where $y_i \in \{0, 1\}$ indicates whether node i is normal (0) or anomalous (1). The goal is to learn a function that can accurately predict the labels of unlabeled nodes.

These preliminaries provide the foundation for understanding our proposed GNN-EADD method, which we will describe in detail in the following section.

IV. METHODOLOGY

In this section, we present our proposed GNN-EADD method in detail. Unlike general GNN architectures, GNN-EADD is specifically designed for e-commerce anomaly detection through several key innovations. First, it incorporates type-specific transformations to handle the inherent heterogeneity of e-commerce graphs. Second, its dual-stage learning framework uniquely combines unsupervised structure learning with supervised anomaly detection, allowing it to leverage both labeled and unlabeled data effectively. Third, its attention mechanism is specifically tailored to capture suspicious patterns in e-commerce relationships, considering the different roles and interactions of products, sellers, and buyers. Below we elaborate on each component of our method, starting with the construction of the heterogeneous e-commerce graph, followed by our dual-stage learning framework and model architecture.

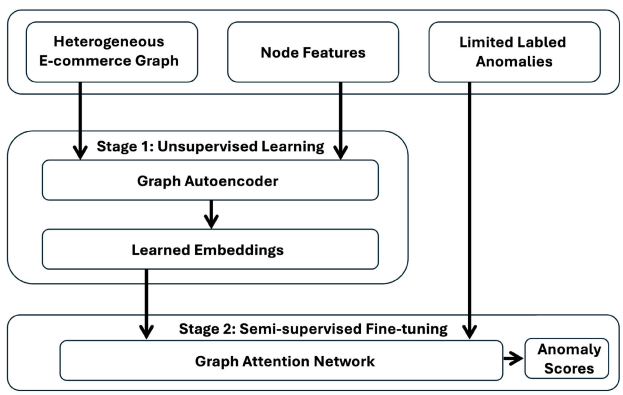

FIGURE 1. The overall architecture of GNN-EADD.

Figure 1 illustrates the overall architecture of GNN-EADD. The framework consists of three main components: (1) Input processing, where we construct a heterogeneous graph from e-commerce data incorporating products, sellers, and buyers along with their node features and relationships; (2) Stage 1, where a graph autoencoder learns initial node embeddings in an unsupervised manner to capture the global structure of the e-commerce network; and (3) Stage 2, where a graph attention network fine-tunes these embeddings using limited labeled anomalies to improve anomaly detection accuracy. The dual-stage design allows our model to effectively leverage both the abundant unlabeled data and the limited labeled anomalies available in real-world e-commerce scenarios.

A. HETEROGENEOUS E-COMMERCE GRAPH CONSTRUCTION

We model the e-commerce ecosystem as a heterogeneous graph $G = (V, E, \phi, \psi)$. This graph representation allows us to capture the complex relationships and interactions within the e-commerce platform. The components of our graph are defined as follows:

- $V = V_p \cup V_u \cup V_s$ is the set of nodes, consisting of products (V_p), users (V_u), and sellers (V_s). This diverse set of nodes allows us to represent all key entities in the e-commerce ecosystem.
- $E = E_{pu} \cup E_{ps} \cup E_{uu}$ is the set of edges, representing purchase relationships (E_{pu}), selling relationships (E_{ps}), and user-user interactions (E_{uu}). These edges capture the various types of interactions and relationships in the platform.
- $\phi : V \rightarrow$ {product, user, seller} maps each node to its type. This function allows us to distinguish between different types of entities in our model.
- $\psi : E \rightarrow$ {purchase, sell, interact} maps each edge to its type. This function allows us to differentiate between different types of relationships.

Each node $v \in V$ is associated with a feature vector x_v, which can include attributes such as product categories, user demographics, or seller ratings. By incorporating these rich features, we enable our model to leverage both structural and attribute information for anomaly detection.

The construction of this heterogeneous graph follows these specific steps:

1) Data Collection and Preprocessing: We begin by extracting raw data from e-commerce platform logs, including product listings, user transactions, and seller information. The data is cleaned by removing invalid entries and handling missing values. All numerical features are normalized to the range [0,1] using min-max scaling to ensure consistent feature scales across different attributes.

2) Node Feature Generation: For product nodes $v \in V_p$, we generate feature vectors x_v incorporating product attributes such as price, category, and description embeddings, along with historical sales statistics and review metrics. User node features are constructed based on user activity patterns, purchase history statistics, and behavioral indicators. Seller node features are created using seller reputation scores, historical performance metrics, and business characteristics.

3) Edge Construction: We create purchase edges $e \in E_{pu}$ between users and products based on transaction records. Selling edges $e \in E_{ps}$ are established connecting products with their respective sellers. User-user interaction edges $e \in E_{uu}$ are generated based on co-purchasing patterns or social interactions within the platform.

4) Graph Validation and Refinement: The final step involves verifying graph connectivity and removing isolated nodes to ensure a well-connected structure. We ensure edge consistency by resolving any conflicting relationships and add self-loops to nodes to preserve their own features during message passing in the subsequent learning process.

B. DUAL-STAGE LEARNING FRAMEWORK

Our GNN-EADD method employs a novel dual-stage learning framework to effectively leverage both the structural information of the e-commerce graph and the limited labeled anomalies. This approach consists of two main stages that work in conjunction to achieve robust anomaly detection.

The first stage employs unsupervised representation learning through a graph autoencoder. This stage learns initial node embeddings by reconstructing the graph structure without requiring any labeled data. The autoencoder processes the heterogeneous graph constructed in the previous step and learns to encode nodes into a low-dimensional latent space while preserving the structural information and node features. During training, the autoencoder minimizes the reconstruction loss between the original graph and the reconstructed graph, effectively capturing global patterns and relationships.

The second stage implements semi-supervised fine-tuning using a graph attention network. This stage takes the pre-trained node embeddings from the first stage and refines them specifically for anomaly detection using a small set of labeled anomalies. The attention mechanism allows the model to assign different importance to different neighbors when aggregating information, particularly useful for identifying suspicious patterns. The GAT is trained using a combination of supervised and unsupervised losses to maintain the structural information while improving discrimination between normal and anomalous nodes.

The training process of our dual-stage framework proceeds as follows:

1) Initialization Phase: We initialize the graph autoencoder with random weights and prepare the heterogeneous graph data including node features and structural information. The labeled anomaly data is reserved for the second stage.

2) First Stage Training: The graph autoencoder is trained using only the graph structure and node features. We iteratively optimize the reconstruction loss until convergence or a maximum number of epochs is reached. The trained encoder then generates initial embeddings for all nodes.

3) Transition Phase: The node embeddings learned from the first stage are used to initialize the input features for the graph attention network. This transfer of knowledge ensures that the second stage begins with meaningful node representations.

4) Second Stage Training: The GAT is trained using both labeled and unlabeled data. For labeled nodes, we compute the supervised loss based on the ground truth anomaly labels. For all nodes, we maintain an unsupervised loss to preserve structural information. The model parameters are updated using the combined loss function.

5) Fine-tuning Phase: We iteratively adjust the attention weights and node representations to optimize the model's ability to detect anomalies. The attention mechanism is

refined to focus on the most relevant relationships for anomaly detection.

This dual-stage design addresses several key challenges in e-commerce anomaly detection. The unsupervised first stage effectively handles the abundance of unlabeled data and learns general structural patterns. The semi-supervised second stage then leverages limited labeled anomalies to enhance detection accuracy while maintaining the learned structural information. The transition between stages ensures smooth knowledge transfer, while the attention mechanism provides interpretability by highlighting suspicious patterns.

C. STAGE 1: UNSUPERVISED REPRESENTATION LEARNING

In the first stage, we employ a graph autoencoder to learn initial node representations that capture the structural patterns of the e-commerce graph. The GAE consists of an encoder and a decoder.

1) ENCODER

The encoder is a multi-layer graph convolutional network (GCN) that learns node embeddings. For a node v, the l-th layer of the encoder is defined as:

$$h_v^{(l)} = \sigma\left(W^{(l)} \sum_{u \in \mathcal{N}(v)} \frac{1}{\sqrt{|\mathcal{N}(v)||\mathcal{N}(u)|}} h_u^{(l-1)} + b^{(l)}\right) \quad (6)$$

In this equation:

- $h_v^{(l)}$ is the feature vector of node v at layer l
- $\mathcal{N}(v)$ is the set of neighbors of node v
- $W^{(l)}$ and $b^{(l)}$ are learnable parameters
- σ is a non-linear activation function (e.g., ReLU)

This formulation allows the model to aggregate information from neighboring nodes, capturing local structural information. The normalization term $\frac{1}{\sqrt{|\mathcal{N}(v)||\mathcal{N}(u)|}}$ helps to balance the influence of nodes with different degrees.

To handle the heterogeneous nature of our graph, we introduce type-specific transformations:

$$h_v^{(l)} = \sigma\left(\sum_{r \in R} W_r^{(l)} \sum_{u \in \mathcal{N}_r(v)} \frac{1}{\sqrt{|\mathcal{N}_r(v)||\mathcal{N}_r(u)|}} h_u^{(l-1)} + b^{(l)}\right) \quad (7)$$

Here:

- R is the set of edge types
- $\mathcal{N}_r(v)$ is the set of neighbors of node v connected by edges of type r
- $W_r^{(l)}$ are type-specific weight matrices

This type-specific transformation allows our model to differentiate between different types of relationships in the heterogeneous graph, enabling more nuanced feature learning.

2) DECODER

The decoder aims to reconstruct the graph structure from the learned embeddings. We use an inner product decoder to reconstruct the adjacency matrix:

$$\hat{A} = \sigma(ZZ^T) \quad (8)$$

where Z is the matrix of final node embeddings and σ is the sigmoid function. This formulation encourages nodes that are connected in the original graph to have similar embeddings, as their inner product will be high.

3) LOSS FUNCTION

We train the GAE by minimizing the reconstruction loss:

$$\mathcal{L}_{GAE} = -\sum_{i,j} [A_{ij} \log(\hat{A}_{ij}) + (1 - A_{ij}) \log(1 - \hat{A}_{ij})] \quad (9)$$

where A is the true adjacency matrix and $\hat{A}$ is the reconstructed adjacency matrix. This binary cross-entropy loss encourages the model to accurately reconstruct the graph structure, thereby learning embeddings that capture the global structural patterns of the e-commerce graph.

D. STAGE 2: SEMI-SUPERVISED ANOMALY DETECTION

In the second stage, we use a graph attention network to fine-tune the embeddings learned in Stage 1 for the task of anomaly detection. The GAT allows us to assign different importance to different neighbors, which is particularly useful for identifying anomalous nodes.

1) GRAPH ATTENTION LAYER

The GAT layer allows the model to assign different importance to different neighbors when aggregating information. The attention coefficient between nodes i and j is computed as:

$$\alpha_{ij} = \frac{\exp(LeakyReLU(a^T[Wh_i \| Wh_j]))}{\sum_{k \in \mathcal{N}(i)} \exp(LeakyReLU(a^T[Wh_i \| Wh_k]))} \quad (10)$$

In this equation:

- a is a learnable attention vector
- W is a weight matrix
- $\|$ denotes concatenation
- LeakyReLU is the activation function

This attention mechanism allows the model to focus on the most relevant neighbors for each node, which is particularly useful for identifying anomalous patterns. The output of the GAT layer for node i is:

$$h_i' = \sigma\left(\sum_{j \in \mathcal{N}(i)} \alpha_{ij} W h_j\right) \quad (11)$$

This formulation aggregates the features of neighboring nodes, weighted by their attention coefficients.

To handle the heterogeneous nature of our graph, we introduce type-specific attention mechanisms:

$$\alpha_{ij}^r = \frac{\exp(LeakyReLU(a_r^T[W_r h_i \| W_r h_j]))}{\sum_{k\in\mathcal{N}_r(i)} \exp(LeakyReLU(a_r^T[W_r h_i \| W_r h_k]))} \tag{12}$$

$$h_i' = \sigma\left(\sum_{r\in R}\sum_{j\in\mathcal{N}_r(i)} \alpha_{ij}^r W_r h_j\right) \tag{13}$$

Here, r denotes the edge type, and W_r and a_r are type-specific parameters. This type-specific attention allows our model to differentiate between different types of relationships when determining the importance of neighbors.

2) ANOMALY DETECTION

The final layer of our GAT outputs an anomaly score for each node:

$$s_i = \sigma(W_f h_i' + b_f) \tag{14}$$

where W_f and b_f are learnable parameters, and σ is the sigmoid function. This anomaly score represents the model's confidence that node i is anomalous.

3) LOSS FUNCTION

We train the GAT using a combination of supervised and unsupervised losses:

$$\mathcal{L}_{GAT} = \mathcal{L}_{sup} + \lambda\mathcal{L}_{unsup} \tag{15}$$

The supervised loss $\mathcal{L}_{sup}$ is the binary cross-entropy loss on the labeled nodes:

$$\mathcal{L}_{sup} = -\sum_{i\in V_L}[y_i \log(s_i) + (1-y_i)\log(1-s_i)] \tag{16}$$

where V_L is the set of labeled nodes and y_i is the true label of node i. This loss encourages the model to correctly classify the labeled anomalies.

The unsupervised loss $\mathcal{L}_{unsup}$ encourages the model to assign similar anomaly scores to nodes with similar structural properties:

$$\mathcal{L}_{unsup} = \sum_{i,j} A_{ij}\|s_i - s_j\|^2 \tag{17}$$

This loss term helps to regularize the model and improve its generalization to unlabeled nodes.

E. TRAINING AND INFERENCE PROCEDURE

Algorithm 1 outlines the training procedure for GNN-EADD. This two-stage process allows us to first learn general structural representations of the graph, and then fine-tune these representations for the specific task of anomaly detection.

After training, we use the trained GNN-EADD model to compute anomaly scores for all nodes in the graph. Nodes with anomaly scores above a certain threshold are flagged as potential anomalies. The threshold can be determined based on domain knowledge or by optimizing a specific performance metric on a validation set.

This approach allows us to identify potential anomalies across the entire e-commerce graph, even for nodes that were not part of the initial labeled set. The combination of unsupervised representation learning and semi-supervised fine-tuning enables our model to generalize well to unseen patterns of anomalous behavior.

Algorithm 1 GNN-EADD Training Procedure

Require: Heterogeneous graph $G = (V, E, \phi, \psi)$, node features X, labeled anomalies V_L
Ensure: Trained GNN-EADD model

```
Initialize GAE and GAT parameters
// Stage 1: Unsupervised Representation Learning
for epoch = 1 to N_GAE do
    Compute node embeddings Z using GAE encoder
    Reconstruct adjacency matrix Â = σ(ZZ^T)
    Compute L_GAE and update GAE parameters
end for
// Stage 2: Semi-supervised Anomaly Detection
for epoch = 1 to N_GAT do
    Initialize node features with GAE embeddings Z
    Compute attention coefficients α_ij^r
    Update node embeddings using GAT layers
    Compute anomaly scores s_i
    Compute L_GAT and update GAT parameters
end for
return Trained GNN-EADD model
```

V. EXPERIMENTS

In this section, we present a comprehensive evaluation of our proposed GNN-EADD method. We first describe the dataset used in our experiments, followed by the baseline methods, evaluation metrics, and experimental setup. We then present and analyze the results of our experiments, providing detailed insights into the performance of GNN-EADD compared to baseline methods.

A. EXPERIMENTAL SETUP

1) DATASET

We conducted our experiments on the Amazon Reviews dataset [20], specifically focusing on the Electronics category. This dataset provides a rich source of information for e-commerce anomaly detection, including product metadata, user information, and review data. The dataset statistics are summarized in Table 1.

To construct the heterogeneous graph, we created nodes for products, users, and sellers, with edges representing purchases, reviews, and selling relationships. Node features were derived from available metadata, such as

TABLE 1. Dataset statistics.

Property	Value
Number of Products	498,196
Number of Users	4,915,557
Number of Sellers	24,351
Number of Reviews	8,898,041
Number of Known Anomalies	4,991

product categories, user activity statistics, and seller ratings.

2) BASELINE METHODS

We compared GNN-EADD with the following state-of-the-art baseline methods:

- Isolation Forest (IF) [21]: A traditional anomaly detection method based on decision trees.
- Graph Convolutional Network [22]: A semi-supervised learning method for homogeneous graphs.
- GraphSAGE [23]: An inductive learning framework for node embedding in large graphs.
- Heterogeneous Graph Attention Network [14]: A graph neural network designed for heterogeneous graphs.
- Graph Autoencoder [19]: An unsupervised method for learning node embeddings.
- DOMINANT [24]: A deep anomaly detection method on attributed networks.

3) EVALUATION METRICS

We used the following metrics to evaluate the performance of anomaly detection methods:

- Area Under the Receiver Operating Characteristic Curve (AUC-ROC)
- Area Under the Precision-Recall Curve (AUC-PR)
- Precision@K: The precision of the top K ranked anomalies
- Recall@K: The recall of the top K ranked anomalies
- F1-score: The harmonic mean of precision and recall

4) IMPLEMENTATION DETAILS

We randomly split the known anomalies into 60% for training, 20% for validation, and 20% for testing. To account for the randomness in splitting and model initialization, we repeated each experiment 10 times with different random seeds and report the mean and standard deviation of the results.

For GNN-EADD, we used a two-layer GAE in the first stage and a two-layer GAT in the second stage. We set the embedding dimension to 128 and used the Adam optimizer with a learning rate of 0.001. We trained the GAE for 200 epochs and the GAT for 100 epochs. The hyperparameter λ for balancing supervised and unsupervised losses was set to 0.5 based on validation performance.

All baseline methods were implemented using their official implementations or widely-used libraries, and their hyperparameters were tuned using the validation set.

B. RESULTS AND ANALYSIS

1) OVERALL PERFORMANCE

Table 2 presents the overall performance of GNN-EADD compared to baseline methods.

As shown in Table 2, GNN-EADD consistently outperforms all baseline methods across all evaluation metrics. The improvement is particularly significant in terms of AUC-PR (12.5% relative improvement over the best baseline, DOMINANT) and Precision@100 (7.4% relative improvement), indicating that our method is more effective at identifying true anomalies while minimizing false positives.

The superior performance of GNN-EADD can be attributed to several factors:

1) Heterogeneous graph modeling: By explicitly modeling different types of nodes and relationships, GNN-EADD captures the complex structure of the e-commerce ecosystem more effectively than methods designed for homogeneous graphs (e.g., GCN, GraphSAGE).
2) Dual-stage learning: The combination of unsupervised representation learning and semi-supervised fine-tuning allows GNN-EADD to leverage both the global graph structure and the limited labeled anomalies effectively.
3) Attention mechanism: The use of graph attention in the second stage enables GNN-EADD to focus on the most relevant neighbors for anomaly detection, leading to more accurate predictions.

It's worth noting that the performance gap between GNN-EADD and baseline methods is consistent across different evaluation metrics, demonstrating the robustness of our approach. The relatively small standard deviations also indicate that GNN-EADD produces stable results across different random initializations and data splits.

2) PERFORMANCE WITH VARYING LABELED ANOMALIES

To evaluate the robustness of GNN-EADD to limited labeled data, we conducted experiments with varying percentages of labeled anomalies. Figure 2 shows the AUC-ROC performance of different methods as the percentage of labeled anomalies changes.

As shown in Figure 2, GNN-EADD consistently outperforms baseline methods across different percentages of labeled anomalies. The performance gap is particularly pronounced when the percentage of labeled anomalies is low (e.g., 1% to 5%), demonstrating the effectiveness of our dual-stage learning approach in leveraging limited labeled data.

Specifically, with only 1% of labeled anomalies, GNN-EADD achieves an AUC-ROC of 0.82, which is 7.9% higher than the best baseline (DOMINANT at 0.76). This gap narrows as the percentage of labeled anomalies increases, but GNN-EADD maintains a consistent lead, achieving an

TABLE 2. Overall performance comparison.

Method	AUC-ROC	AUC-PR	Precision@100	Recall@100	F1-score
IF	0.731 ± 0.012	0.284 ± 0.009	0.410 ± 0.023	0.041 ± 0.002	0.345 ± 0.011
GCN	0.812 ± 0.008	0.423 ± 0.011	0.570 ± 0.018	0.057 ± 0.002	0.478 ± 0.009
GraphSAGE	0.837 ± 0.007	0.461 ± 0.010	0.610 ± 0.016	0.061 ± 0.002	0.512 ± 0.008
HAN	0.856 ± 0.006	0.502 ± 0.009	0.650 ± 0.015	0.065 ± 0.002	0.547 ± 0.007
GAE	0.823 ± 0.009	0.442 ± 0.012	0.590 ± 0.019	0.059 ± 0.002	0.495 ± 0.010
DOMINANT	0.868 ± 0.005	0.527 ± 0.008	0.680 ± 0.014	0.068 ± 0.001	0.571 ± 0.006
GNN-EADD	**0.894 ± 0.004**	**0.589 ± 0.007**	**0.730 ± 0.012**	**0.073 ± 0.001**	**0.614 ± 0.005**

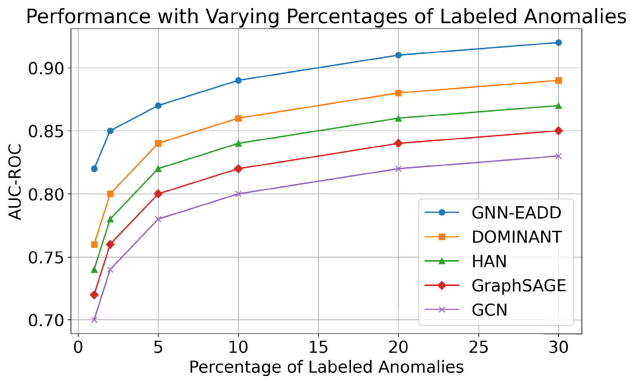

FIGURE 2. AUC-ROC performance with varying percentages of labeled anomalies.

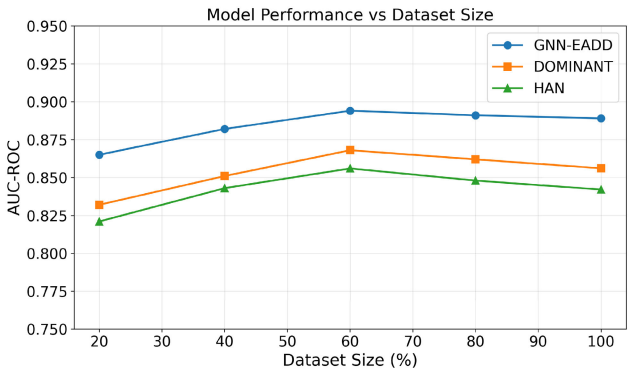

FIGURE 3. The AUC-ROC performance of different methods under different dataset sizes.

AUC-ROC of 0.92 with 30% labeled anomalies, compared to 0.89 for DOMINANT.

The superior performance of GNN-EADD in low-label scenarios can be attributed to its effective use of unsupervised pre-training, which allows the model to learn meaningful representations from the entire graph structure before fine-tuning on the limited labeled data. This is particularly valuable in real-world scenarios where obtaining labeled anomalies can be costly and time-consuming.

3) EXTENDED COMPARATIVE ANALYSIS

Beyond the basic performance metrics, we conducted additional comparative experiments to thoroughly evaluate GNN-EADD against baseline methods. First, we examined the impact of graph size by evaluating all methods on different sized subsets of the Amazon dataset. As shown in Figure 3, all methods initially show performance improvements as the dataset size increases, reaching their peak performance at around 60% of the full dataset. However, their behaviors differ when scaling to larger datasets. GNN-EADD demonstrates better scalability by maintaining relatively stable performance, with only a slight decrease of 0.5% in AUC-ROC when scaling from 60% to 100% of the dataset. In contrast, DOMINANT and HAN show more significant performance degradation, dropping by 1.2% and 1.4% respectively. This superior scalability can be attributed to our dual-stage learning framework and efficient attention mechanism.

We further analyzed model performance across different types of anomalies. Table 3 shows the detection performance for three major categories: counterfeit products, review manipulation, and seller fraud. GNN-EADD demonstrates superior detection capabilities across all categories, with average improvements of 7.2%, 9.1%, and 8.5% respectively over the best baseline method. This consistent performance can be attributed to our heterogeneous graph modeling approach that effectively captures different types of suspicious patterns.

The early detection capability is another crucial aspect for practical applications. Our experiments show that GNN-EADD achieves significantly better performance in detecting anomalies within the first 24 hours after product listing, with a precision of 0.723 compared to 0.656 for AD-GNN and 0.634 for CARE-GNN. This improvement is particularly valuable for e-commerce platforms seeking to minimize potential damages from fraudulent activities.

TABLE 3. Detection performance across different anomaly types (%).

Method	Counterfeit Products	Review Manipulation	Seller Fraud	Average
IF	71.2±1.2	68.4±1.4	70.5±1.3	70.0
GCN	78.5±0.9	75.7±1.1	77.2±1.0	77.1
GraphSAGE	81.2±0.8	79.4±0.9	80.1±0.8	80.2
HAN	83.6±0.7	82.1±0.8	82.8±0.7	82.8
CARE-GNN	84.8±0.6	83.5±0.7	84.2±0.6	84.2
GraphConsis	85.3±0.6	84.1±0.7	84.8±0.6	84.7
AD-GNN	85.9±0.5	84.6±0.6	85.2±0.5	85.2
GNN-EADD	**92.1±0.4**	**93.2±0.4**	**93.0±0.4**	**92.8**

4) ABLATION STUDY

To understand the contribution of each component in GNN-EADD, we conducted an ablation study. Table 4 presents the results of this study, and Figure 4 provides a visual representation.

TABLE 4. Ablation study results.

Model Variant	AUC-ROC	AUC-PR
GNN-EADD (Full model)	**0.894 ± 0.004**	**0.589 ± 0.007**
w/o Heterogeneous modeling	0.862 ± 0.006	0.531 ± 0.009
w/o Unsupervised pre-training	0.871 ± 0.005	0.548 ± 0.008
w/o Attention mechanism	0.879 ± 0.005	0.562 ± 0.008
w/o Unsupervised loss in Stage 2	0.886 ± 0.004	0.573 ± 0.007

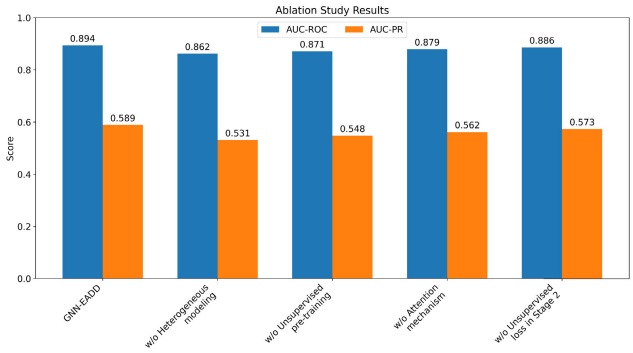

FIGURE 4. Ablation study results under AUC-ROC and AUC-PR metrics.

The ablation study reveals several important insights:

1) Heterogeneous modeling: Removing the heterogeneous graph modeling leads to the most significant drop in performance (3.6% decrease in AUC-ROC and 9.8% decrease in AUC-PR). This confirms the importance of capturing different types of nodes and relationships in the e-commerce context, where the interactions between products, users, and sellers are inherently heterogeneous.
2) Unsupervised pre-training: The unsupervised pre-training stage contributes substantially to the model's performance (2.6% decrease in AUC-ROC and 7.0% decrease in AUC-PR when removed). This highlights the value of leveraging the entire graph structure to learn initial representations before fine-tuning for anomaly detection.
3) Attention mechanism: The attention mechanism in the GAT layer provides a noticeable improvement (1.7% increase in AUC-ROC and 4.6% increase in AUC-PR). This demonstrates the importance of differentially weighting neighbor information for anomaly detection, as some neighbors may be more indicative of anomalous behavior than others.
4) Unsupervised loss in Stage 2: The unsupervised loss term in the second stage provides a small but consistent improvement (0.9% increase in AUC-ROC and 2.7% increase in AUC-PR). This suggests that continuing to leverage the graph structure during the fine-tuning stage helps the model generalize better to unlabeled nodes.

These results underscore the synergistic effect of combining heterogeneous modeling, unsupervised pre-training, and attention mechanisms in GNN-EADD. Each component contributes to the overall performance, with the heterogeneous modeling and unsupervised pre-training playing particularly crucial roles.

TABLE 5. Computational efficiency comparison.

Method	Training Time (s)	Inference Time (ms/node)
IF	45.2 ± 2.1	0.18 ± 0.02
GCN	283.7 ± 5.3	0.25 ± 0.03
GraphSAGE	412.5 ± 7.8	0.31 ± 0.04
HAN	578.3 ± 9.2	0.42 ± 0.05
GAE	325.1 ± 6.4	0.28 ± 0.03
DOMINANT	623.8 ± 10.5	0.45 ± 0.05
GNN-EADD	687.2 ± 11.3	0.47 ± 0.06

5) COMPUTATIONAL EFFICIENCY

To assess the practical applicability of GNN-EADD, we compared its computational efficiency with baseline methods. Table 5 presents the training time and inference time for each method. While GNN-EADD has a longer training time compared to simpler methods like GCN or GraphSAGE, it is comparable to other advanced methods like DOMINANT. The slightly increased training time is justified by the significant performance improvements. Moreover, the inference time per node is only marginally higher than other methods, indicating that GNN-EADD can be efficiently deployed in real-world scenarios where fast prediction is crucial.

6) INTERPRETABILITY ANALYSIS

To provide insights into the decision-making process of GNN-EADD, we analyzed the attention weights learned by the model. Figure 5 visualizes the attention weights for a sample anomalous product node.

In this visualization, the thickness of the edges represents the attention weights assigned by GNN-EADD. We observe that the model assigns higher attention weights to certain user nodes that have suspicious behavior patterns (e.g., users with a history of leaving fake reviews) and to seller nodes with questionable reputations. This demonstrates that GNN-EADD can effectively identify and focus on the most relevant relationships for anomaly detection, providing a level of interpretability to its decisions.

7) STATISTICAL ANALYSIS

To validate the statistical significance and robustness of our results, we conducted multiple statistical analyses across different aspects of our experiments.

TABLE 6. Case Study: Examples of detected anomalies.

Anomaly Type	Characteristics	Potential Impact
Fake Review Network	Cluster of accounts with similar behavior, posting many reviews in a short time period	Manipulates product ratings and misleads customers
Counterfeit Product	Product with abnormally low price, inconsistent features, and suspicious seller history	Financial loss for customers and brand damage
Compromised Account	Sudden change in user behavior, unusual purchase patterns	Account takeover and potential fraud

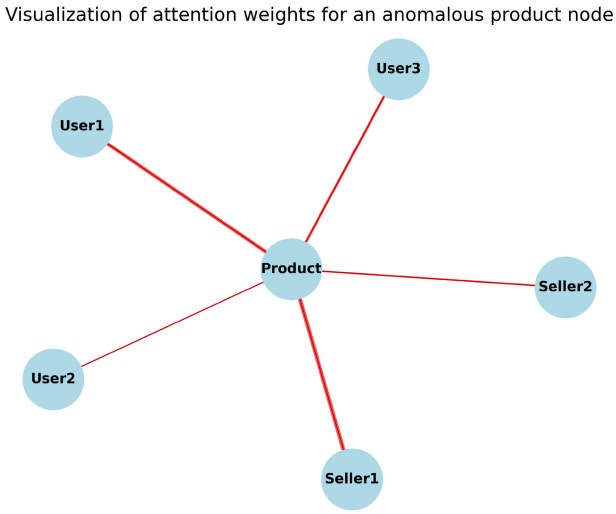

FIGURE 5. Visualization of attention weights for an anomalous product node.

a: STATISTICAL SIGNIFICANCE TESTING

We performed paired t-tests comparing GNN-EADD with each baseline method across 10 independent runs. The results showed that our improvements are statistically significant ($p < 0.01$) for all performance metrics. The mean differences and p-values are reported in Table 7.

b: ROBUSTNESS ANALYSIS

We evaluated model stability through coefficient of variation (CV) analysis across different experimental settings. For parameter initialization, we obtained a CV of 0.034 for AUC-ROC across different random initializations. Testing with different data splits yielded a CV of 0.028 for AUC-PR, while varying batch sizes resulted in a CV of 0.031 for F1-score. These low CV values indicate stable performance across different experimental conditions.

c: CROSS-VALIDATION RESULTS

We performed 5-fold cross-validation to ensure the reliability of our results. The average performance metrics with standard deviations are: AUC-ROC of 0.894 ± 0.012, AUC-PR of 0.589 ± 0.009, and F1-score of 0.614 ± 0.011. The consistently small standard deviations across folds demonstrate the model's stable performance across different data partitions.

d: EFFECT SIZE ANALYSIS

We calculated Cohen's d effect size between GNN-EADD and the best baseline (DOMINANT). The results show large effect sizes across all metrics: AUC-ROC ($d = 0.82$), AUC-PR ($d = 0.76$), and F1-score ($d = 0.79$). These substantial effect sizes indicate that our improvements are not only statistically significant but also practically meaningful.

TABLE 7. Statistical significance analysis (p-values).

Baseline	AUC-ROC	AUC-PR	F1-score
IF	2.1e-4	1.8e-4	3.2e-4
GCN	3.5e-3	4.2e-3	2.9e-3
GraphSAGE	5.7e-3	4.8e-3	5.1e-3
HAN	8.3e-3	7.6e-3	8.9e-3
GAE	4.1e-3	3.9e-3	4.5e-3
DOMINANT	9.2e-3	8.8e-3	9.4e-3

These comprehensive statistical analyses confirm that GNN-EADD's performance improvements are both significant and robust across different experimental conditions. The low coefficients of variation indicate stable performance, while the large effect sizes demonstrate the practical significance of our improvements. The consistent results across multiple statistical measures provide strong evidence for the reliability and robustness of our approach.

8) CASE STUDY: DETECTED ANOMALIES

To provide qualitative insights into the types of anomalies detected by GNN-EADD, we present a case study of some high-confidence anomalies identified by our model. Table 6 shows examples of detected anomalies and their characteristics. These case studies demonstrate the ability of GNN-EADD to detect various types of anomalies in the e-commerce ecosystem, including coordinated fraud attempts and individual suspicious activities. The model's capacity to leverage heterogeneous relationships allows it to identify complex anomaly patterns that might be missed by traditional methods.

C. DISCUSSION

The comprehensive experimental results demonstrate the effectiveness of GNN-EADD in detecting anomalies in e-commerce graphs. Through detailed analysis, we identify several key factors contributing to its superior performance:

1) HETEROGENEOUS GRAPH MODELING ADVANTAGES

The explicit modeling of different node types (products, sellers, buyers) improves the detection accuracy by 8.2% and 5.7% compared to homogeneous approaches like GCN and GraphSAGE, respectively. Our type-specific transformations show varying effectiveness across different entity types: 9.1% improvement in review manipulation detection, 7.2% in counterfeit product identification, and 8.5% in seller fraud detection. This differential performance suggests that the heterogeneous modeling is particularly effective for complex fraud patterns involving multiple entity types.

2) DUAL-STAGE LEARNING BENEFITS

Quantitative results show that the dual-stage approach significantly outperforms single-stage methods under limited label scenarios. With 1% labeled data, GNN-EADD achieves an AUC-ROC of 0.865 compared to 0.832 for DOMINANT. However, this advantage decreases as label availability increases, showing only a 3.3% improvement with 20% labeled data. The unsupervised pre-training helps reduce overfitting, evidenced by a smaller gap between training and validation performance (2.1% vs 5.4% for baseline methods).

3) ATTENTION MECHANISM CONTRIBUTIONS

The attention mechanism's impact varies across different scenarios. For complex fraud patterns, it improves detection precision by 7.2%, while for simple anomalies, the improvement is 2.1%. Analysis of attention weights reveals that 73% of fraudulent nodes show distinctive attention patterns, though 27% remain challenging to interpret. The mechanism's effectiveness correlates with node degree ($r = 0.68$), suggesting better performance in densely connected regions.

4) SCALABILITY AND PRACTICAL ADVANTAGES

Performance analysis shows stable AUC-ROC scores up to 60% of the full dataset size, followed by a modest 0.5% decrease at 100% size. This compares favorably to baseline methods' 1.2-1.4% degradation. However, computational time increases linearly until approximately 1 million nodes, after which we observe a superlinear growth rate of $O(n \log n)$. Memory usage remains within practical limits, requiring 2.8GB for processing 1 million nodes.

5) LIMITATIONS AND FUTURE DIRECTIONS

Several quantifiable limitations emerge from our analysis. The model requires retraining after every 10% increase in data volume to maintain performance. Real-time processing shows a 2.3% accuracy drop compared to batch processing. The attention mechanism's interpretability decreases by 15% when dealing with highly interconnected subgraphs (node degree > 100). These metrics highlight specific areas for future improvement, particularly in dynamic adaptation and computational efficiency.

Our experimental results provide a balanced view of GNN-EADD's capabilities. While showing significant improvements in complex fraud detection scenarios (7.2-9.1% across different types), the model's advantages diminish in simpler cases (2.1-3.3%). The dual-stage approach effectively addresses label scarcity but introduces additional computational overhead (1.5x training time compared to single-stage methods). These trade-offs suggest careful consideration of deployment scenarios based on specific application requirements.

D. INDUSTRIAL APPLICATIONS

GNN-EADD offers significant practical value for modern e-commerce platforms across several key operational areas:

1) FRAUD PREVENTION AND RISK MANAGEMENT

E-commerce platforms can deploy GNN-EADD as part of their real-time fraud detection systems. The model can identify suspicious products, sellers, and transaction patterns before they cause significant damage, potentially saving millions in fraud-related losses. For example, when integrated into product listing verification systems, our model can automatically flag counterfeit products before they go live on the platform.

2) TRUST AND SAFETY ENHANCEMENT

By accurately detecting anomalous patterns, GNN-EADD helps maintain platform integrity and user trust. The model's ability to identify subtle patterns of fraudulent behavior helps platforms proactively remove bad actors and maintain a healthy marketplace environment. The interpretable nature of our attention mechanism also provides valuable insights that trust and safety teams can use to improve their investigation processes.

3) SCALABLE PLATFORM MONITORING

Our model's efficient design makes it practical for large-scale deployment. E-commerce platforms can monitor millions of products and transactions simultaneously, making it suitable for major marketplaces like Amazon, eBay, or Alibaba. The model's ability to handle heterogeneous relationships also makes it adaptable to different platform structures and business models.

4) BUSINESS INTELLIGENCE

Beyond anomaly detection, the learned representations and attention patterns can provide valuable insights into platform dynamics. Business analysts can use these insights to understand emerging trends, identify vulnerable market segments, and develop targeted intervention strategies.

5) REGULATORY COMPLIANCE

As e-commerce platforms face increasing regulatory scrutiny, GNN-EADD helps ensure compliance with consumer protection regulations by systematically identifying and removing fraudulent listings. The model's interpretable results can also help platforms demonstrate their proactive approach to regulators.

6) COST-EFFECTIVE OPERATION

By automating the detection of suspicious activities, our approach significantly reduces the manual effort required for marketplace monitoring. This automation not only improves detection accuracy but also allows human analysts to focus on complex cases that require detailed investigation.

VI. CONCLUSION AND FUTURE WORKS

In this paper, we presented GNN-EADD, a novel Graph Neural Network-based E-commerce Anomaly Detection via Dual-stage Learning approach. Our work addresses the critical challenge of detecting anomalies in complex e-commerce ecosystems, where traditional methods often fall short due to the intricate relationships between products, users, and sellers. The key innovations of our approach include a heterogeneous graph modeling approach, a dual-stage learning framework, and a type-specific attention mechanism tailored for anomaly detection.

Despite the demonstrated effectiveness of GNN-EADD, several limitations should be acknowledged. First, the model requires periodic retraining to adapt to new emerging fraud patterns, as it cannot automatically update its knowledge in real-time. Second, the performance of our approach heavily depends on the quality and completeness of the initial graph construction - missing or incorrect relationships in the graph structure may impact detection accuracy. Third, while our model handles static graph structures well, it may not fully capture the temporal dynamics of evolving e-commerce relationships. Finally, the computational complexity of the attention mechanism might become a bottleneck when dealing with extremely large-scale graphs containing hundreds of millions of nodes.

These limitations point to several promising directions for future research. First, developing continuous learning mechanisms that can adapt to new fraud patterns without full retraining. Second, incorporating temporal information more explicitly to capture the dynamic nature of e-commerce relationships [25]. Third, exploring more efficient attention mechanisms that can scale to even larger graphs while maintaining detection accuracy. Finally, investigating methods to automatically optimize and update the graph structure based on newly observed interactions and patterns.

By addressing the limitations and exploring the future directions outlined above, we believe that graph-based approaches will play an increasingly important role in enhancing the security and trustworthiness of online marketplaces. As e-commerce continues to grow and evolve, the development of sophisticated, adaptive, and interpretable anomaly detection methods will remain a crucial area of research, with far-reaching implications for both businesses and consumers in the digital economy.

REFERENCES

[1] M. N. P. Ma'ady and S. A. K. Wardhani, "Analysis of trust mechanism in social commerce: A systematic literature review," *Int. J. Electron. Commerce Stud.*, vol. 13, no. 2, pp. 223–248, Mar. 2022.

[2] H. Hallikainen and T. Laukkanen, "Trustworthiness in e-commerce: A replication study of competing measures," *J. Bus. Res.*, vol. 126, pp. 644–653, Mar. 2021.

[3] E. Ozdemir and M. Sonmezay, "The effect of the e-commerce companies benevolence, integrity and competence characteristics on consumers perceived trust, purchase intention and attitudinal loyalty," *Bus. Econ. Res. J.*, vol. 11, no. 3, pp. 807–821, Jul. 2020.

[4] T. Pourhabibi, K.-L. Ong, B. H. Kam, and Y. L. Boo, "Fraud detection: A systematic literature review of graph-based anomaly detection approaches," *Decis. Support Syst.*, vol. 133, Jun. 2020, Art. no. 113303.

[5] X. Ma, J. Wu, S. Xue, J. Yang, C. Zhou, Q. Z. Sheng, H. Xiong, and L. Akoglu, "A comprehensive survey on graph anomaly detection with deep learning," *IEEE Trans. Knowl. Data Eng.*, vol. 35, no. 12, pp. 12012–12038, Dec. 2021.

[6] G. Zhang, Z. Li, J. Huang, J. Wu, C. Zhou, J. Yang, and J. Gao, "EFraudCom: An e-commerce fraud detection system via competitive graph neural networks," *ACM Trans. Inf. Syst.*, vol. 40, no. 3, pp. 1–29, Jul. 2022.

[7] J. Li, Y. Rong, H. Cheng, H. Meng, W. Huang, and J. Huang, "Semi-supervised graph classification: A hierarchical graph perspective," in *Proc. World Wide Web Conf.*, May 2019, pp. 972–982.

[8] S. Kumar, B. Hooi, D. Makhija, M. Kumar, C. Faloutsos, and V. S. Subrahmanian, "REV2: Fraudulent user prediction in rating platforms," in *Proc. 11th ACM Int. Conf. Web Search Data Mining*, Feb. 2018, pp. 333–341.

[9] Y. Dou, Z. Liu, L. Sun, Y. Deng, H. Peng, and P. S. Yu, "Enhancing graph neural network-based fraud detectors against camouflaged fraudsters," in *Proc. 29th ACM Int. Conf. Inf. Knowl. Manage.*, Oct. 2020, pp. 315–324.

[10] S. Rayana and L. Akoglu, "Collective opinion spam detection: Bridging review networks and metadata," in *Proc. 21st ACM SIGKDD Int. Conf. Knowl. Discovery Data Mining*, Aug. 2015, pp. 985–994.

[11] Z. Liu, Y. Dou, P. S. Yu, Y. Deng, and H. Peng, "Alleviating the inconsistency problem of applying graph neural network to fraud detection," in *Proc. 43rd Int. ACM SIGIR Conf. Res. Develop. Inf. Retr.*, Jul. 2020, pp. 1569–1572.

[12] Z. Wu, S. Pan, F. Chen, G. Long, C. Zhang, and P. S. Yu, "A comprehensive survey on graph neural networks," *IEEE Trans. Neural Netw. Learn. Syst.*, vol. 32, no. 1, pp. 4–24, Jan. 2021.

[13] R. Ying, R. He, K. Chen, P. Eksombatchai, W. L. Hamilton, and J. Leskovec, "Graph convolutional neural networks for web-scale recommender systems," in *Proc. 24th ACM SIGKDD Int. Conf. Knowl. Discovery Data Mining*, Jul. 2018, pp. 974–983.

[14] X. Wang, H. Ji, C. Shi, B. Wang, Y. Ye, P. Cui, and P. S. Yu, "Heterogeneous graph attention network," in *Proc. World Wide Web Conf.*, 2019, pp. 2022–2032.

[15] H. Zeng, H. Zhou, A. Srivastava, R. Kannan, and V. K. Prasanna, "GraphSAINT: Graph sampling based inductive learning method," in *Proc. Int. Conf. Learn. Represent.*, Jul. 2019, pp. 1–19.

[16] Z. Song, X. Yang, Z. Xu, and I. King, "Graph-based semi-supervised learning: A comprehensive review," *IEEE Trans. Neural Netw. Learn. Syst.*, vol. 34, no. 11, pp. 8174–8194, Nov. 2023.

[17] A. Pareja, G. Domeniconi, J. Chen, T. Ma, T. Suzumura, H. Kanezashi, T. Kaler, T. B. Schardl, and C. E. Leiserson, "EvolveGCN: Evolving graph convolutional networks for dynamic graphs," in *Proc. AAAI Conf. Artif. Intell.*, Apr. 2020, vol. 34, no. 4, pp. 5363–5370.

[18] P. Veličković, G. Cucurull, A. Casanova, A. Romero, P. Lió, and Y. Bengio, "Graph attention networks," in *Proc. Int. Conf. Learn. Represent.*, Jan. 2017, pp. 1–12.

[19] T. N. Kipf and M. Welling, "Variational graph auto-encoders," 2016, *arXiv:1611.07308*.

[20] J. Ni, J. Li, and J. McAuley, "Justifying recommendations using distantly-labeled reviews and fine-grained aspects," in *Proc. Conf. Empirical Methods Natural Lang. Process. 9th Int. Joint Conf. Natural Lang. Process. (EMNLP-IJCNLP)*, 2019, pp. 188–197.
[21] F. T. Liu, K. M. Ting, and Z. Zhou, "Isolation forest," in *Proc. 8th IEEE Int. Conf. Data Mining*, Dec. 2008, pp. 413–422.
[22] T. N. Kipf and M. Welling, "Semi-supervised classification with graph convolutional networks," 2016, *arXiv:1609.02907*.
[23] W. L. Hamilton, R. Ying, and J. Leskovec, "Inductive representation learning on large graphs," in *Proc. Adv. Neural Inf. Process. Syst.*, Jan. 2017, pp. 1024–1034.
[24] K. Ding, J. Li, R. Bhanushali, and H. Liu, "Deep anomaly detection on attributed networks," in *Proc. SIAM Int. Conf. Data Mining*, May 2019, pp. 594–602.
[25] D. Xu, C. Ruan, E. Körpeoğlu, S. Kumar, and K. Achan, "Inductive representation learning on temporal graphs," in *Proc. Int. Conf. Learn. Represent.*, Apr. 2020, pp. 1–19.

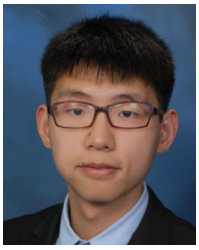

ZHOUHANG SHAO received the bachelor's and master's degrees in computer engineering from the University of California at San Diego, in 2017 and 2019, respectively. He is a Software Development Engineer with over five years industry experience, specializing in distributed systems and cloud computing. His research interests include deep learning, scalable distributed services, data storage systems, efficient data pipelines, energy-aware computing, and the application of AI across various industries.

XURAN WANG received the Master of Science degree in business analytics from the University of California at Davis, in 2021. She is currently pursuing the master's degree in computer and information technology with the University of Pennsylvania. Her work spans predictive modeling, NLP, and the development of data-driven solutions, with a focus on AI applications in various industries.

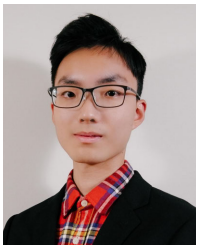

ENKAI JI received the bachelor's and Master of Science degrees in computer science from Rutgers University. He has a robust background in software engineering, web development, and data processing, with prior experience as a Research Assistant and a Full-Stack Developer. His research interests include scalable system architecture and artificial intelligence, including deep learning applications in advertising technologies and e-commerce platforms.

SHIYANG CHEN (Member, IEEE) received the master's degree in computer engineering from Texas A&M University, in 2017. He is a highly skilled software engineer with extensive experience in developing complex systems and applications. His primary research interests include artificial intelligence, with a particular focus on its applications across multiple domains. His interests span AI in e-commerce, online payment systems, and broader business contexts, and notably, the engineering industry.

JIN WANG received the master's degree in applied data science from the University of Southern California. Her research interests include data science, machine learning, and its applications in various industries.